

EN3160: Image Processing and Machine Vision

Assignment 1: Intensity Transformations and Neighborhood Filtering

Department of Electronic and Telecommunication Engineering

University of Moratuwa

September 14, 2026

Name: Samarasinghe S.M.R.R.

Index Number: 230566U

GitHub Repository: github.com/RumethRanhinda/EN3160-Vision-UoM

1. Intensity Transformation Implementation

```
[ ]: import cv2
import numpy as np
import matplotlib.pyplot as plt

def intensity_transform(im, breakpoints):
    bp = np.array(breakpoints)
    xp = bp[:, 0] # Input intensities
    yp = bp[:, 1] # Output intensities

    # Apply interpolation
    transformed_im = np.interp(im.flatten(), xp, yp).reshape(im.shape)

    return transformed_im.astype(np.uint8)

img = cv2.imread('Images/emma.jpeg', cv2.IMREAD_GRAYSCALE)
bp_given = [[0, 0], [49, 49], [50, 100], [149, 255], [150, 150], [255, 255]] # Provided ↵
↵breakpoints
bp_custom = [[0, 0], [50, 20], [175, 100], [255, 255]] # Custom breakpoints

transformed_given = intensity_transform(img, bp_given)
transformed_custom = intensity_transform(img, bp_custom)

# Grid ticks for the plots
grid_ticks = [0, 50, 100, 150, 200, 255]





fig, axes = plt.subplots(2, 3, figsize=(15, 9))
# Given Breakpoints
bp_given_arr = np.array(bp_given)
axes[0, 0].plot(bp_given_arr[:,0], bp_given_arr[:,1], marker='o', color='#55499C')
axes[0, 0].set_title("Transformation Plot (Given)")
axes[0, 0].set_xlabel("Input intensity"); axes[0, 0].set_ylabel("Output intensity")
axes[0, 0].set_xlim([0, 255]); axes[0, 0].set_ylim([0, 260])
# Apply custom grid
axes[0, 0].set_xticks(grid_ticks); axes[0, 0].set_yticks(grid_ticks); axes[0, 0].grid(True)
axes[0, 1].imshow(img, cmap='gray', vmin=0, vmax=255)
axes[0, 1].set_title("Original Image"); axes[0, 1].axis('off')
```

```

axes[0, 2].imshow(transformed_given, cmap='gray', vmin=0, vmax=255)
axes[0, 2].set_title("Transformed Image (Given)"); axes[0, 2].axis('off')

# Custom Breakpoints
bp_custom_arr = np.array(bp_custom)
axes[1, 0].plot(bp_custom_arr[:,0], bp_custom_arr[:,1], marker='o', color='#55499C')
axes[1, 0].set_title("Transformation Plot (Custom)")
axes[1, 0].set_xlabel("Input intensity"); axes[1, 0].set_ylabel("Output intensity")
axes[1, 0].set_xlim([0, 255]); axes[1, 0].set_ylim([0, 260])
axes[1, 0].set_xticks(grid_ticks); axes[1, 0].set_yticks(grid_ticks); axes[1, 0].grid(True)
axes[1, 1].imshow(img, cmap='gray', vmin=0, vmax=255)
axes[1, 1].set_title("Original Image"); axes[1, 1].axis('off')
axes[1, 2].imshow(transformed_custom, cmap='gray', vmin=0, vmax=255)
axes[1, 2].set_title("Transformed Image (Custom)"); axes[1, 2].axis('off')
plt.tight_layout(); plt.show()

```



The given intensity transformation made high intensities above 150 darker in a discontinuous way which explains the artifact seen in the output image.

The custom transformation was implemented with the intention of highlighting the freckles within the bright area of the face by providing a wider range for brighter pixels.

2. Accentuate White and Gray Matter

```

[3]: brain_img = cv2.imread('Images/brain.jpeg', cv2.IMREAD_GRAYSCALE)

if brain_img is not None:
    bp_white_matter = [[0, 0], [150, 5], [255, 255]]
    bp_gray_matter = [[0, 0], [128, 0], [190, 255], [255, 255]]
    wm_accentuated = intensity_transform(brain_img, bp_white_matter)

```



```

gm_accentuuated = intensity_transform(brain_img, bp_gray_matter)
grid_ticks = [0, 50, 100, 150, 200, 255]
fig, axes = plt.subplots(2, 2, figsize=(10, 10))

# White Matter Transformation
bp_wm_arr = np.array(bp_white_matter)
axes[0, 0].plot(bp_wm_arr[:,0], bp_wm_arr[:,1], marker='o', color='#55499C')
axes[0, 0].set_title("Transformation: White Matter")
axes[0, 0].set_xlabel("Input intensity"); axes[0, 0].set_ylabel("Output intensity")
axes[0, 0].set_xlim([0, 255]); axes[0, 0].set_ylim([-5, 260])
axes[0, 0].set_xticks(grid_ticks); axes[0, 0].set_yticks(grid_ticks); axes[0, 0].grid(True)
axes[0, 1].imshow(wm_accentuuated, cmap='gray', vmin=0, vmax=255)
axes[0, 1].set_title("White Matter Accentuated"); axes[0, 1].axis('off')

# Gray Matter Transformation
bp_gm_arr = np.array(bp_gray_matter)
axes[1, 0].plot(bp_gm_arr[:,0], bp_gm_arr[:,1], marker='o', color='#55499C')
axes[1, 0].set_title("Transformation: Gray Matter"); axes[1, 0].set_xlabel("Input intensity")
axes[1, 0].set_ylabel("Output intensity"); axes[1, 0].set_xlim([0, 255]); axes[1, 0].
↪set_ylim([-5, 260])
axes[1, 0].set_xticks(grid_ticks); axes[1, 0].set_yticks(grid_ticks); axes[1, 0].grid(True)
axes[1, 1].imshow(gm_accentuuated, cmap='gray', vmin=0, vmax=255)
axes[1, 1].set_title("Gray Matter Accentuated"); axes[1, 1].axis('off')
plt.tight_layout()
plt.show()
else:
    print("Image not found")

```



3. Gamma Correction

```
[4]: img_bgr = cv2.imread('Images/legs.jpg')

if img_bgr is not None:
    # For display in Matplotlib
    img_rgb = cv2.cvtColor(img_bgr, cv2.COLOR_BGR2RGB)

    # Convert to L*a*b* color space
    img_lab = cv2.cvtColor(img_bgr, cv2.COLOR_BGR2LAB)
    L, a, b = cv2.split(img_lab)
```

```

gamma = 0.5

# Normalize L to [0, 1], apply power law, and scale back to [0, 255]
L_normalized = L / 255.0
L_corrected = np.power(L_normalized, gamma)
L_corrected = np.uint8(L_corrected * 255)

# Merge the corrected L plane with the original a and b planes
img_lab_corrected = cv2.merge((L_corrected, a, b))
img_corrected_rgb = cv2.cvtColor(img_lab_corrected, cv2.COLOR_LAB2RGB)

# Compute histograms
hist_L_orig = cv2.calcHist([L], [0], None, [256], [0, 256])
hist_L_corr = cv2.calcHist([L_corrected], [0], None, [256], [0, 256])

fig, axes = plt.subplots(2, 2, figsize=(12, 10))

# Original Image
axes[0, 0].imshow(img_rgb); axes[0, 0].set_title("Original Image"); axes[0, 0].axis('off')


# Gamma Corrected Image
axes[0, 1].imshow(img_corrected_rgb); axes[0, 1].set_title(rf"Gamma Corrected Image_{\gamma={gamma}}")
↪ axes[0, 1].axis('off')

# Histogram of Original L Plane
axes[1, 0].plot(hist_L_orig, color='#55499C')
axes[1, 0].set_title("Histogram: Original $L$ Plane")
axes[1, 0].set_xlabel("Pixel Intensity")
axes[1, 0].set_ylabel("Frequency")
axes[1, 0].set_xlim([0, 255])
axes[1, 0].grid(True)

# Histogram of Corrected L Plane
axes[1, 1].plot(hist_L_corr, color='#55499C')
axes[1, 1].set_title("Histogram: Corrected $L$ Plane")
axes[1, 1].set_xlabel("Pixel Intensity")
axes[1, 1].set_ylabel("Frequency")
axes[1, 1].set_xlim([0, 255])
axes[1, 1].grid(True)

plt.tight_layout()
plt.show()
else:
    print("Image not found")

```



4. Enhancing Vibrance using Intensity Transformation

```
[5]: img_bgr = cv2.imread('Images/spider.jpeg')

if img_bgr is not None:
    img_rgb = cv2.cvtColor(img_bgr, cv2.COLOR_BGR2RGB)

    # Convert to HSV color space
    img_hsv = cv2.cvtColor(img_bgr, cv2.COLOR_BGR2HSV)
    H, S, V = cv2.split(img_hsv)

    a = 0.55
    sigma = 70.0

    # Convert S plane to float32 to prevent overflow during mathematical operations
    S_float = S.astype(np.float32)

    gaussian_term = np.exp(-((S_float - 128.0)**2) / (2.0 * sigma**2))
    S_transformed = S_float + a * 128.0 * gaussian_term
    S_transformed = np.clip(S_transformed, 0, 255).astype(np.uint8)

    # Recombine the three planes
```

```

img_hsv_enhanced = cv2.merge((H, S_transformed, V))
img_enhanced_rgb = cv2.cvtColor(img_hsv_enhanced, cv2.COLOR_HSV2RGB)

# Generate the 1D intensity transformation curve
x_vals = np.arange(0, 256, dtype=np.float32)
y_vals = x_vals + a * 128.0 * np.exp(-((x_vals - 128.0)**2) / (2.0 * sigma**2))
y_vals = np.clip(y_vals, 0, 255)

fig, axes = plt.subplots(1, 3, figsize=(15, 5))



# Plot 1: Transformation Curve
axes[0].plot(x_vals, y_vals, color='#55499C', linewidth=2)
axes[0].plot([0, 255], [0, 255], 'k--', alpha=0.5) # Reference linear line
axes[0].set_title(rf"Saturation Transformation ($a={a}$)")
axes[0].set_xlabel("Input Saturation ($x$)")
axes[0].set_ylabel("Output Saturation ($f(x)$)")
axes[0].set_xlim([0, 255]); axes[0].set_ylim([0, 255])
axes[0].grid(True)

# Plot 2: Original Image
axes[1].imshow(img_rgb)
axes[1].set_title("Original Image")
axes[1].axis('off')

# Plot 3: Enhanced Image
axes[2].imshow(img_enhanced_rgb)
axes[2].set_title("Vibrance-Enhanced Image")
axes[2].axis('off')

plt.tight_layout()
plt.show()
else:
    print("Image not found")

```



Increasing a improves the vibrancy of the image. Ranging a from 0.5 to 0.8 did not result in significant changes. Although above 0.8 the image starts to become unrealistically colourful. To avoid major clipping of saturation values at 255, a is chosen to be 0.55.

5. Custom Histogram Equalization

```
[ ]: def histogram_equalization(img):
    # Compute the histogram
    hist, bins = np.histogram(img.flatten(), 256, [0, 256])
    cdf = hist.cumsum()

    # 3. Mask out zero values in the CDF to ignore empty bins during normalization
    cdf_masked = np.ma.masked_equal(cdf, 0)

    # 4. Normalize the CDF to the range [0, 255]
    cdf_normalized = (cdf_masked - cdf_masked.min()) * 255 / (cdf_masked.max() - cdf_masked.
    min())

    # 5. Fill the masked values back with 0 and cast to uint8
    cdf_final = np.ma.filled(cdf_normalized, 0).astype(np.uint8)

    # 6. Map the original pixel intensities to the new equalized values
    img_eq = cdf_final[img]

    return img_eq, hist, cdf_final

img = cv2.imread('Images/shells.tif', cv2.IMREAD_GRAYSCALE)

if img is not None:
    img_equalized, hist_orig, mapping_function = histogram_equalization(img)

    # Compute the histogram
    hist_eq, _ = np.histogram(img_equalized.flatten(), 256, [0, 256])
    fig, axes = plt.subplots(2, 2, figsize=(12, 10))

    # Original Image
    axes[0, 0].imshow(img, cmap='gray', vmin=0, vmax=255)
    axes[0, 0].set_title("Original Image")
    axes[0, 0].axis('off')

    # Original Histogram
    axes[0, 1].plot(hist_orig, color='#55499C')
    axes[0, 1].fill_between(range(256), hist_orig, color='#55499C', alpha=0.3)
    axes[0, 1].set_title("Histogram before Equalization")
    axes[0, 1].set_xlabel("Pixel Intensity")
    axes[0, 1].set_ylabel("Frequency")
    axes[0, 1].set_xlim([0, 255])
    axes[0, 1].grid(True)


    # Equalized Image
    axes[1, 0].imshow(img_equalized, cmap='gray', vmin=0, vmax=255)
    axes[1, 0].set_title("Equalized Image")
    axes[1, 0].axis('off')

    # Equalized Histogram
    axes[1, 1].plot(hist_eq, color='#55499C')
    axes[1, 1].fill_between(range(256), hist_eq, color='#55499C', alpha=0.3)
    axes[1, 1].set_title("Histogram after Equalization")
    axes[1, 1].set_xlabel("Pixel Intensity")
    axes[1, 1].set_ylabel("Frequency")
    axes[1, 1].set_xlim([0, 255])
    axes[1, 1].grid(True)


    plt.tight_layout()
```

```
plt.show()
else:
    print("Image not found")
```

Original Image




Histogram before Equalization



Equalized Image



Histogram after Equalization



6. Neighborhood Filtering

```
[ ]: img_bgr = cv2.imread('Images/jennifer.jpeg')
img_rgb = cv2.cvtColor(img_bgr, cv2.COLOR_BGR2RGB)

# convert to HSV color space and split into H, S, V planes
img_hsv = cv2.cvtColor(img_bgr, cv2.COLOR_BGR2HSV)
H, S, V = cv2.split(img_hsv)

# Threshold mask on saturation plane / Otsu's method calculates the threshold value.
_, mask = cv2.threshold(S, 0, 255, cv2.THRESH_BINARY + cv2.THRESH_OTSU)

# Obtain the foreground
fg_v = cv2.bitwise_and(V, V, mask=mask)
```

```

# Calculate the histogram of the Value plane, strictly using the mask to ignore the background
hist_fg = cv2.calcHist([V], [0], mask, [256], [0, 256]).flatten()

# (d) Obtain the cumulative sum of the histogram using np.cumsum
cdf = np.cumsum(hist_fg)

# (e) Histogram-equalize the foreground[cite: 1]
# Mask out 0s to prevent empty bins from skewing the normalization
cdf_m = np.ma.masked_equal(cdf, 0)
cdf_normalized = (cdf_m - cdf_m.min()) * 255 / (cdf_m.max() - cdf_m.min())
cdf_final = np.ma.filled(cdf_normalized, 0).astype(np.uint8)

# Create a copy of the V plane and map the equalized values ONLY to the foreground coordinates
V_eq = V.copy()
fg_coords = mask > 0
V_eq[fg_coords] = cdf_final[V[fg_coords]]

# (f) Extract the background and add with the histogram equalized foreground[cite: 1]
mask_inv = cv2.bitwise_not(mask)
bg_v = cv2.bitwise_and(V, V, mask=mask_inv)
fg_eq_v = cv2.bitwise_and(V_eq, V_eq, mask=mask)

# Add the untouched background and the equalized foreground back together
final_v = cv2.add(bg_v, fg_eq_v)

# Recombine the H, S, and new V planes, and convert back to RGB for display
final_hsv = cv2.merge((H, S, final_v))
final_rgb = cv2.cvtColor(final_hsv, cv2.COLOR_HSV2RGB)

# Display the required planes and outputs[cite: 1]
fig, axes = plt.subplots(2, 3, figsize=(15, 10))


# Row 1: The individual HSV planes in grayscale
axes[0, 0].imshow(H, cmap='gray'); axes[0, 0].set_title('Hue Plane')
axes[0, 1].imshow(S, cmap='gray'); axes[0, 1].set_title('Saturation Plane')
axes[0, 2].imshow(V, cmap='gray'); axes[0, 2].set_title('Value Plane')

# Row 2: Mask, Original, and Final Result
axes[1, 0].imshow(mask, cmap='gray'); axes[1, 0].set_title('Foreground Mask')
axes[1, 1].imshow(img_rgb); axes[1, 1].set_title('Original Image')
axes[1, 2].imshow(final_rgb); axes[1, 2].set_title('Equalized Foreground Image')

for ax in axes.flatten():
    ax.axis('off')

plt.tight_layout()
plt.show()

```



The saturation plane is able to be used for the foreground mask because the background is dull and thus have lower saturation. The thresholding do miss a portion of the black dress the subject is wearing, which can be ignored.